

KE YANG

Department of Computer Science
The University of Texas at San Antonio
San Pedro I #410 E,
506 Dolorosa St,
San Antonio, TX 78204

ke.yang@utsa.edu
[CAREAI @ UTSA](#)
[Personal Website](#)
[Google Scholar Profile](#)
+1 (215) 900-5854

CURRENT AFFILIATIONS

Assistant Professor

Department of Computer Science,
College of Sciences,
The University of Texas at San Antonio

Aug, 2023 - current

Founding Faculty

Cohort for AI REsponsibility (CAREAI),
The University of Texas at San Antonio

Sep, 2023 - current

Core Faculty

School of Data Science,
The University of Texas at San Antonio

Sep, 2023 - current

Faculty Member

AI Consortium for Human Well-being (MATRIX),
The University of Texas at San Antonio

Oct, 2023 - current

PAST ACADEMIC APPOINTMENTS

Postdoctoral Research Associate

Manning College of Information & Computer Sciences,
University of Massachusetts Amherst

Sep, 2021 - Aug, 2023

Postdoctoral Fellow

Center for Data Science,
University of Massachusetts Amherst

Sep, 2021 - Aug, 2023

Visiting Scholar

Tandon School of Engineering,
New York University

2021 - 2022

RESEARCH INTEREST

My work is broadly in the areas of *AI trustworthiness, responsibility and safety, data management, machine learning, and human-centered data science*. In particular, I focus on topics such as *hallucination of Large Language Models, explainable AI, and algorithmic accountability*. Other areas of focus include *AI and machine learning education* and public engagement.

EDUCATION

- Ph.D. in Computer Science*, New York University 2019 - 2021
- Advisor: Julia Stoyanovich, GPA 3.9/4.0
- Thesis: [Fairness, diversity, and interpretability in ranking](#)
- Committee: Julia Stoyanovich (chair), Daniel Neill, Oded Nov, and Sebastian Schelter
- Ph.D. candidate in Computer Science*, [Drexel University](#) (transferred) 2015 - 2018
- Advisor: Julia Stoyanovich, GPA 3.98/4.0
- M.S. in Computer Science*, [Beijing Technology and Business University](#), 2012 - 2015
- Advisors: Qian Mo and Zhongming Han, GPA 3.51/4.0
- Thesis: Detecting web spammers in online shopping websites
- Graduated *summa cum laude*
- B.Eng. in Software Engineering*, [Beijing Technology and Business University](#) 2008 - 2012
- Advisor: Qian Mo, GPA 3.45/4.0
- Graduated *summa cum laude*
-

PAST PROFESSIONAL EXPERIENCE

- Postdoctoral Research Associate**
Manning College of Information & Computer Sciences,
University of Massachusetts Amherst, MA, U.S.A. 2021 - 2023
- Supervisor and co-PI: [Alexandra Meliou](#)
- Fully funded by CDS Postdoctoral Fellowship
- Selected publication: [Non-Invasive Fairness in Learning through the Lens of Data Drift, 2023](#)
- Graduate Research Assistant**
[VIDA Research Center](#), New York University, NY, U.S.A. 2019 - 2021

- Supervisor: [Julia Stoyanovich](#)
- Fully funded by a graduate research assistantship
- Selected publication: [Fairness in Ranking: A Survey, 2021](#) (received 145 citations until Aug, 2023)
- More of my past projects on dataresponsibly.github.io I built for the lab during graduate school.

Research Intern

AT&T Labs, New York City, NY, U.S.A.

2019

- Supervisors: [Emily Dodwell](#), [Ritwik Mitra](#), and [Balachander Krishnamurthy](#)
- Project: Fairness and transparency in machine learning

Graduate Research Assistant

College of Computing & Informatics, Drexel University, PA, U.S.A.

2015 - 2018

- Supervisor: [Julia Stoyanovich](#)
- Fully funded by a graduate research assistantship
- Selected publication:
 - > [Measuring Fairness in Ranked Outputs, 2017](#) (received 341 citations until Aug, 2023)
 - > [A Nutritional Label for Rankings, 2018](#) (received 113 citations until Aug, 2023)

Research engineer

Elite & Resource (start-up company), Beijing, China

2014 - 2015

- Supervisor: [Peng Sun](#)
- Project: Preventing flood disasters using machine learning
- Results have been integrated as a core component of a national data management system for floor disasters

Graduate Research Assistant

Beijing Technology and Business University, Beijing, China

2012 - 2015

- Supervisors: [Zhongming Han](#) and [Qian Mo](#)
- Fully funded by a graduate research assistantship
- Selected publication:
 - > [Overview of Web Spammer Detection, 2013](#) (published in a top-tier computer science journal in Chinese)
 - > [Analyzing Spectrum Features of Weight User Relation Graph to Identify Large Spammer Groups in Online Shopping Websites, 2015](#) (published in a top-tier computer science journal in Chinese)

PUBLICATIONS

Peer-reviewed Conference Publications

1. Md Sohanur Rahman, Yuexia Zhang, Anthony Rios, and **Ke Yang**. Extracting Causal Relational Rules for Medical Question-Answering Tasks using Large Language Models. Workshop on Generative AI for Smart and Connected Health: Innovations, Challenges, and Applications (GenAI4SCH), co-located with IEEE/ACM conference on Connected Health: Applications, Systems and Engi-

neering Technologies (**CHASE 2025**), New York, USA, regular short paper.

2. Hilal Pataci, **Ke Yang**, and Anthony Rios. Empowering Privacy Policy Understanding and Accessibility in Smart Home Devices. **INFORMS 2024**, presented as a poster.
3. Youmi Suk and **Ke Yang**. Using Conditional Tabular Generative Adversarial Networks for process data generation in Monte Carlo simulation studies. Paper presented at the International Meeting of Psychometric Society (**IMPS 2024**), Prague, Czech Republic.
4. **Ke Yang** and Alexandra Meliou. Non-Invasive Fairness in Learning through the Lens of Data Drift. IEEE International Conference on Data Engineering (**ICDE 2024**), Utrecht, Netherlands.
5. **Ke Yang**. DEEILS: Data Ethics Embedded Interactive Learning System for Computer Science Students. The Technical Symposium on Computer Science Education (**SIGCSE TS 2024**), Portland, Oregon, United States, presenting as a lighting talk for junior faculty.
6. **Ke Yang**, Joshua R. Loftus, and Julia Stoyanovich. Causal Intersectionality and Fair Ranking. Symposium on Foundations of Responsible Computing, (**FORC 2021**), Virtual Conference.
7. Meike Zehlike, **Ke Yang**, and Julia Stoyanovich. Mapping Normative Frameworks to Fair Ranking. ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization (**EAAMO 2021**), presented as a poster.
8. **Ke Yang**, Vasilis Gkatzelis, and Julia Stoyanovich. Balanced Ranking with Diversity Constraints. International Joint Conference on Artificial Intelligence (**IJCAI 2019**), Macao, China.
9. **Ke Yang**, Julia Stoyanovich, Abolfazl Asudeh, Bill Howe, H.V. Jagadish, and Gerome Miklau. A Nutritional Label for Rankings. ACM International Conference on Management of Data (**SIGMOD 2018 Demonstration**), Houston, USA.
10. Julia Stoyanovich, **Ke Yang**, H.V. Jagadish. Online Set Selection with Fairness and Diversity Constraints. International Conference on Extending Database Technology (**EDBT 2018**), Vienna, Austria.
11. **Ke Yang** and Julia Stoyanovich. Measuring Fairness in Ranked Outputs. International Conference on Scientific and Statistical Database Management (**SSDBM 2017**), Chicago, USA.

Peer-reviewed Journal Publications

12. Zehlike Meike, **Ke Yang**, and Julia Stoyanovich. Fairness in Ranking, Part I: Score-based Ranking. ACM Computing Surveys (**CSUR 2022**).
13. Zehlike Meike, **Ke Yang**, and Julia Stoyanovich. Fairness in Ranking, Part II: Learning-to-Rank and Recommender Systems. ACM Computing Surveys (**CSUR 2022**).
14. Zhongming Han, **Ke Yang**, and Xusheng Tan. Analyzing Spectrum Features of Weight User Relation Graph to Identify Large Spammer Groups in Online Shopping Websites. **Chinese Journal of Computers**, 40(4): 939-954 (2017) (in Chinese).
15. Zhongming Han, **Ke Yang**, Fengmin Xu, and Dagao Duan. Probabilistic Graphical Model for Detecting Spammers in Microblog Websites. **Journal of Embedded Systems**, 8(1): 12-23 (2016).
16. Qian Mo and **Ke Yang**. Overview of Web Spammer Detection. **Chinese Journal of Software**, 25(7): 1505-1526 (2014) (in Chinese).

Peer-reviewed Workshops

17. **Ke Yang**, Biao Huang, Julia Stoyanovich, and Sebastian Schelter. Fairness-Aware Instrumentation of Preprocessing Pipelines for Machine Learning. Workshop on Human-In-the-Loop Data Analytics in conjunction with the 2020 ACM SIGMOD/PODS Conference, (**HILDA@SIGMOD 2020**), Portland, OR, USA.

Under review

- * Youmi Suk, Chenguang Pan, and **Ke Yang**. Using Generative AI for Sequential Data Generation in Monte Carlo Simulation Studies, 2024.

GRANTS, FELLOWSHIPS, AWARDS, & HONORS

- Reducing Hallucination of LLMs for Causality-integrated Personal Medical History Question Answering, Collaborative Seed Funding Grant (\$35,000), School of Data Science (SDS) and the Open Cloud Institute (OCI), UTSA, 2024
- Faculty Travel Support Program, Internal funding, Office of Research, UTSA, April 2024
- Sponsored participant for [NSF Secure and Trustworthy Cyberspace program \(SaTC\) Aspiring PI Workshop](#), Chicago, May 2024
- Sponsored faculty for [NCFDD Faculty Success Program summer 2024](#), Office of Research, UTSA
- UTSA funding for research equipment (\$250,000), [STARs Program](#), the University of Texas System, 2023
- Postdoctoral Fellowship, *Center for Data Science*, University of Massachusetts Amherst, 2021
- Pearl Brownstein Doctoral Research Award, *Tandon School of Engineering*, New York University, 2021
- Invited researcher, [Fairness in Machine Learning workshop at Google](#), 2018

TEACHING

Instructor at the University of Texas at San Antonio

- CS 5823 *Trust, Confidence, and Explainability in AI*, Spring 2025
- CS 7123 *Research Methods*, Spring 2024
- CS 5443 *Database Management Systems*, Fall 2023, 2024

Assisting Instructor at University of Massachusetts Amherst

- CMPSCI 345 *Practice and Applications of Data Management*, Fall 2022

Teaching Assistant at Drexel University

- CS 500 *Fundamentals of Databases*, Winter 2017
 - CS 461 *Database Systems*, Spring 2017
-

STUDENTS' MENTORING

Current

- Sinmi Akindele, Ph.D. student, University of Texas at San Antonio, College of Sciences, joined in Fall 2023
- Sohanur Rahman, Ph.D. student, University of Texas at San Antonio, College of Sciences, joined in Fall 2023
- Asif Ahmed, Ph.D. student, University of Texas at San Antonio, College of Sciences, joined in fall 2023
- Kevin Mathew Maran, Master's student, University of Texas at San Antonio, College of Sciences, joined in fall 2024
- Sai Siri Chittineni, undergraduate student, University of Texas at San Antonio, Computer Science, joined in fall 2023
- Anudeep Bonagiri, high school student, Tom C Clark High School and Northwest Vista College, joined in spring 2025

Past

- Jason Johnson, Master's student, University of Texas at San Antonio, College of Sciences, participated in spring and fall 2024
 - Samyuktha Arcot Megha Raj, Master's student, University of Texas at San Antonio, College of Sciences, participated in fall 2024
 - Ashmit Sethi, high school student, Dublin High School, participated in fall 2024
 - Rajarshi Ghosh, high school student, Oak High School and Lone Star College, participated during spring 2024
 - Zafeiria Moumoulidou, Ph.D. student, University of Massachusetts Amherst, Manning College of Information & Computer Sciences, co-advised with Prof. Alexandra Meliou in 2022 and 2023
 - Alexandra Burushkina, Master's student, University of Massachusetts Amherst, Center for Data Science, [Data Science for Social Good Summer Program](#), graduated June 2023, current position: Ph.D. student in Computer Science at University of Massachusetts Amherst
 - Peter Tran, Master's student, University of Massachusetts Amherst, Center for Data Science, [Data Science for Social Good Summer Program](#), graduated June 2023, current position: Ph.D. student in Psychometrics at University of Massachusetts Amherst
 - Pragath Siva, high school student, Sharon High School, graduated June 2023, current position: freshman in Computer Science and Economics at the University of Illinois Urbana-Champaign
 - Samasth Norway Ananda, Master's student, New York University, graduated June 2021, current position: software developer at Oracle
 - Biao Huang, Master's student, New York University, graduated June 2020, current position: software developer in China lead tech company
 - Jensine Raihan, undergraduate student, New York University, 2020
-

INVITED TALKS

- Data-driven Solutions for AI Responsibility, [Virtual Seminar](#), [DATA Lab @ Northeastern University](#), 2024 Spring
- AI and the Changing Nature of Work, Panel Speaker, [South Texas Human Resources Symposium \(STXHRS\) 2024](#)

- Causal Intersectionality and Fair Ranking, [Integrated Symposium on Algorithmic and Testing Fairness in Education: Methods and Applications](#), SREE 2023 Conference Embedding Equity for Better Methods, Evidence, and Practice, September 2023
 - Mitigating Bias in Data-driven Tasks, [UTSA MATRIX](#), Fall Seminar Series, September 2023
 - Responsible Data Science, Brown University, virtual, March 2023
 - Reducing the Impact of Bias on Data-driven Tasks, Ohio State University, February 2023
 - Data Science for Social Good, Panel Speaker, Central Massachusetts Women in Data Science (WiDS), Worcester Polytechnic Institute (WPI), March 2022
 - Data Bias and their Impact on Machine Learning, [MIT Horizon](#) event, November 2022
 - Fairness and Diversity in Rankings, AT&T Lab Research Weekly Symposium, June 2019
 - Fairness In Ranking Systems, AT&T Lab Graduate Student Symposium, November 2018
-

ACADEMIC SERVICE

Conference Organization

- Session co-chair: Responsibility and Privacy at Extending Database Technology (EDBT) 2022
- Technical chair: 4th Workshop on Data Management for End-to-End Machine Learning (DEEM) at SIGMOD 2020

Program Committee Member

- Guided AI Workshop @ ACM Conference on Management of Data (SIGMOD) 2024
- International Conference on Extending Database Technology (EDBT) Industry Program 2023, 2022
- ACM Web Conference (a.k.a. International World Wide Web Conference, abbreviated as WWW) 2024 Industry, 2023 Industry, Workshops, Tutorials, Posters, and Demos, 2022 Industry
- IEEE Conference on Data Engineering (ICDE) Demo Track, 2023, 2022
- ACM Conference on Fairness, Accountability, and Transparency (ACM FAccT) 2021
- ACM Workshop on Responsible Recommendation (FAccTRec) at RecSys 2024, 2023, 2022, 2021

- ACM Workshop on Data Management for End-to-End Machine Learning (DEEM) at SIGMOD 2021, 2020

Ad-hoc Reviewer

- Ethics reviewer, ICML 2025
- Ethics reviewer, NeurIPS 2024, 2023, 2022, 2021
- ACM Conference on Management of Data (SIGMOD) 2023
- ACM Transactions on Social Computing 2023
- ACM Conference on Human Factors in Computing Systems (CHI) 2021

Other Activities

- Social Science Panelist for American Association of University Women (AAUW), International Fellowships, 2025, 2024
- Student volunteer: ACM Conference on Fairness, Accountability, and Transparency (ACM FAT*) 2019